

1 Suppose that we are working under the Gauss-Markov model

$$\mathbf{Y} = \mathbf{X}\beta + \epsilon$$

where $E(\epsilon) = \mathbf{0}$ and $\text{var}(\epsilon) = \sigma^2 \mathbf{I}$. Let $\hat{\mathbf{Y}}$ be the ordinary least square estimator of \mathbf{Y} .

(a) Show that $\hat{\mathbf{Y}}$ and $\mathbf{Y} - \hat{\mathbf{Y}}$ are uncorrelated.

Let $\hat{\mathbf{Y}} = P_X \mathbf{Y} = X(X'X)^{-1}X'\mathbf{Y}$.

First, note that $\hat{\mathbf{Y}}$ and $\mathbf{Y} - \hat{\mathbf{Y}}$ are orthogonal.

$$\hat{\mathbf{Y}}'(\mathbf{Y} - \hat{\mathbf{Y}}) = (\mathbf{P}_X \mathbf{Y})'(\mathbf{Y} - \mathbf{P}_X \mathbf{Y}) = \mathbf{Y}'\mathbf{P}_X'(\mathbf{Y} - \mathbf{P}_X \mathbf{Y}) = \mathbf{Y}'\mathbf{P}_X \mathbf{Y} - \mathbf{Y}'\mathbf{P}_X' \mathbf{P}_X \mathbf{Y} = \mathbf{Y}'\mathbf{P}_X \mathbf{Y} - \mathbf{Y}'\mathbf{P}_X \mathbf{Y} = \mathbf{0}$$

Therefore, the expectation

$$E(\hat{\mathbf{Y}}'(\mathbf{Y} - \hat{\mathbf{Y}})) = E(\mathbf{0}) = \mathbf{0}$$

$$E(\hat{\mathbf{Y}}) = E(\mathbf{X}\hat{\beta}) = XE(\hat{\beta}) = X\beta$$

$$E(\mathbf{Y} - \hat{\mathbf{Y}}) = X\beta - X\beta = \mathbf{0}$$

This shows $\hat{\mathbf{Y}}$ and $\mathbf{Y} - \hat{\mathbf{Y}}$ are uncorrelated because $E(\hat{\mathbf{Y}}(\mathbf{Y} - \hat{\mathbf{Y}})) - E(\hat{\mathbf{Y}})E(\mathbf{Y} - \hat{\mathbf{Y}})$ is zero.

(b) Show that

$$E\{(\mathbf{Y} - \hat{\mathbf{Y}})^T (\mathbf{Y} - \hat{\mathbf{Y}})\} = \sigma^2 \{n - \text{rank}(\mathbf{X})\}.$$

You may use Theorem 5.2a of R&S.

Theorem 5.2a states: If \mathbf{y} is a random vector with mean μ and covariance matrix Σ and if \mathbf{A} is a symmetric matrix of constants, then

$$E(\mathbf{y}'\mathbf{A}\mathbf{y}) = \text{tr}(\mathbf{A}\Sigma) + \mu'\mathbf{A}\mu.$$

We can apply Theorem 5.2a with $\mathbf{A} = \mathbf{I}$ and $\mathbf{y} = \mathbf{Y} - \hat{\mathbf{Y}}$. From part a) above, our mean μ is $\mathbf{0}$. The term $\mu'\mathbf{A}\mu$ is therefore $\mathbf{0}$, and the expression of interest becomes:

$$E\{(\mathbf{Y} - \hat{\mathbf{Y}})^T (\mathbf{Y} - \hat{\mathbf{Y}})\} = \text{tr}(\Sigma)$$

Where Σ is the covariance matrix of $\mathbf{Y} - \hat{\mathbf{Y}}$.

$$\begin{aligned} \text{Var}(\mathbf{Y} - \hat{\mathbf{Y}}) &= \text{Var}((\mathbf{I} - \mathbf{P}_X)\mathbf{Y}) \\ &= (\mathbf{I} - \mathbf{P}_X)\text{Var}(\mathbf{Y})(\mathbf{I} - \mathbf{P}_X)' \\ &= (\mathbf{I} - \mathbf{P}_X)\text{Var}(\mathbf{Y})(\mathbf{I} - \mathbf{P}_X) \\ &= (\mathbf{I} - \mathbf{P}_X)\sigma^2 \mathbf{I}(\mathbf{I} - \mathbf{P}_X) \\ &= \sigma^2 (\mathbf{I} - \mathbf{P}_X) \end{aligned}$$

Trace of $\mathbf{I} - \mathbf{P}_X$ is $\text{tr}(\mathbf{I}) - \text{tr}(\mathbf{P}_X)$. The trace of an $n \times n$ identity matrix \mathbf{I} is n , and the trace a projection matrix is the rank of target space, $\text{tr}(P_X) = \text{rank}(X)$. The trace of product of a scalar c and a matrix \mathbf{A} is the product $\text{tr}(c\mathbf{A}) = c\text{tr}(\mathbf{A})$. Thus, $\text{tr}(\sigma^2(\mathbf{I} - \mathbf{P}_X)) = \sigma^2(n - \text{rank}(\mathbf{X}))$. This gives the desired result:

$$E\{(\mathbf{Y} - \hat{\mathbf{Y}})^T (\mathbf{Y} - \hat{\mathbf{Y}})\} = \sigma^2 \{n - \text{rank}(\mathbf{X})\}.$$

2 Consider the one-way ANOVA model $y_{ij} = \mu + \tau_i + \epsilon_{ij}$ for the j th individual of the i th group .

Suppose there are 4 treatments (groups) and the sample sizes are respectively 2,1,1,2 for treatments. Now suppose that $\mathbf{Y} = (y_{11}, y_{12}, y_{21}, y_{31}, y_{41}, y_{42})^T = (2, 1, 4, 6, 3, 5)^T$ contains the observations.

Use R and weighted generalized least squares to find an appropriate estimate for

$$E(\mathbf{Y}) \text{ and } \begin{pmatrix} 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 1 \end{pmatrix} \beta$$

in the Aiken model with $\text{var}(\epsilon) = \mathbf{V}$ for two cases where

(a) $\mathbf{V} = \mathbf{V}_1 = \text{diag}(1, 9, 9, 1, 1, 9)$

The full model described in this question in $\mathbf{Y} = \mathbf{X}\beta + \epsilon$ matrix form is:

$$\begin{pmatrix} y_{11} \\ y_{12} \\ y_{21} \\ y_{31} \\ y_{41} \\ y_{42} \end{pmatrix} = \begin{pmatrix} 2 \\ 1 \\ 4 \\ 6 \\ 3 \\ 5 \end{pmatrix} = \begin{pmatrix} 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} \mu \\ \tau_1 \\ \tau_2 \\ \tau_3 \\ \tau_4 \end{pmatrix} + \begin{pmatrix} \epsilon_{11} \\ \epsilon_{12} \\ \epsilon_{21} \\ \epsilon_{31} \\ \epsilon_{41} \\ \epsilon_{42} \end{pmatrix}$$

We have $\text{var}(\epsilon) = \sigma^2 \mathbf{V}$, so we must re-write the model in terms of $\mathbf{U} = \mathbf{V}^{-1/2} \mathbf{Y}$ as follows:

$$\begin{aligned} \mathbf{V} &= \mathbf{V}^{1/2} \mathbf{V}^{1/2}, \mathbf{V} \text{ is a diagonal matrix} \\ \text{Let } \mathbf{U} &= \mathbf{V}^{-1/2} \mathbf{Y} \\ E(\mathbf{U}) &= \mathbf{V}^{-1/2} E\mathbf{Y} = \mathbf{V}^{-1/2} \mathbf{X}\beta \\ &= \mathbf{W}\beta \\ \text{Var}(\mathbf{U}) &= \mathbf{V}^{-1/2} \text{Var}(\mathbf{Y}) \mathbf{V}^{-1/2} \\ &= \sigma^2 \mathbf{V}^{-1/2} \mathbf{V} \mathbf{V}^{-1/2} \\ &= \sigma^2 \mathbf{I} \\ \epsilon^* &= \mathbf{V}^{-1/2} \epsilon \\ E(\epsilon^*) &= E(\mathbf{V}^{-1/2} \epsilon) \\ &= \mathbf{V}^{-1/2} E(\epsilon) \\ &= \mathbf{0} \\ \text{Var}(\epsilon^*) &= \text{Var}(\mathbf{V}^{-1/2} \epsilon) \\ &= \mathbf{V}^{-1/2} \text{Var}(\epsilon) \mathbf{V}^{-1/2} \\ &= \mathbf{V}^{-1/2} \sigma^2 \mathbf{V} \mathbf{V}^{-1/2} \\ &= \sigma^2 \mathbf{I} \end{aligned}$$

This gives us $\mathbf{U} = \mathbf{W}\beta + \epsilon^*$, where the Gauss-Markov assumptions hold for \mathbf{U} .

We first calculate $\mathbf{V}^{-1/2}$:

$$\mathbf{V} = \mathbf{V}_1 = \text{diag}(1, 9, 9, 1, 1, 9)$$

$$\text{So, } \mathbf{V}_1^{1/2} = \text{diag}(1, 3, 3, 1, 1, 3)$$

$$\mathbf{V}_1^{-1/2} = \text{diag}(1, 1/3, 1/3, 1, 1, 1/3)$$

We can check this in R,

```
V <- diag(c(1,9,9,1,1,9))
Vhi <- solve(V^(1/2))
lvector(Vhi)
```

$$\mathbf{V}^{-1/2} = \begin{pmatrix} 1.00 & 0.00 & 0.00 & 0.00 & 0.00 & 0.00 \\ 0.00 & 0.33 & 0.00 & 0.00 & 0.00 & 0.00 \\ 0.00 & 0.00 & 0.33 & 0.00 & 0.00 & 0.00 \\ 0.00 & 0.00 & 0.00 & 1.00 & 0.00 & 0.00 \\ 0.00 & 0.00 & 0.00 & 0.00 & 1.00 & 0.00 \\ 0.00 & 0.00 & 0.00 & 0.00 & 0.00 & 0.33 \end{pmatrix}$$

$$\mathbf{U} = \mathbf{V}^{-1/2} \mathbf{Y} = \begin{pmatrix} y_{11} \\ \frac{1}{3}y_{12} \\ \frac{1}{3}y_{21} \\ y_{31} \\ y_{41} \\ \frac{1}{3}y_{42} \end{pmatrix} = \begin{pmatrix} 2 \\ \frac{1}{3} \\ \frac{4}{3} \\ 6 \\ 3 \\ \frac{5}{3} \end{pmatrix}$$

Checking \mathbf{U} in R gives:

```
Y <- matrix(c(2, 1, 4, 6, 3, 5), nrow=6, ncol=1)
U <- Vhi %*% Y
lvector(U)
```

$$\mathbf{U} = \begin{pmatrix} 2.00 \\ 0.33 \\ 1.33 \\ 6.00 \\ 3.00 \\ 1.67 \end{pmatrix}$$

$$\mathbf{W} = \mathbf{V}^{-1/2} \mathbf{X}$$

$$= \text{diag}(1, 1/3, 1/3, 1, 1, 1/3) \begin{pmatrix} 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 1 \end{pmatrix}$$

$$= \begin{pmatrix} 1 & 1 & 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & 0 & \frac{1}{3} & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 1 \\ \frac{1}{3} & 0 & 0 & 0 & \frac{1}{3} \end{pmatrix}$$

Checking \mathbf{W} in R gives:

```
X <- matrix(c(rep(1,6),
               1,1,0,0,0,0,
               0,0,1,0,0,0,
               0,0,0,1,0,0,
               0,0,0,0,1,1),nrow = 6,byrow=FALSE)
W <- Vhi %*% X
lvector(W)
```

$$\mathbf{W} = \begin{pmatrix} 1.00 & 1.00 & 0.00 & 0.00 & 0.00 \\ 0.33 & 0.33 & 0.00 & 0.00 & 0.00 \\ 0.33 & 0.00 & 0.33 & 0.00 & 0.00 \\ 1.00 & 0.00 & 0.00 & 1.00 & 0.00 \\ 1.00 & 0.00 & 0.00 & 0.00 & 1.00 \\ 0.33 & 0.00 & 0.00 & 0.00 & 0.33 \end{pmatrix}$$

(a).1 Solving for $\widehat{E(\mathbf{Y})}$: $\mathbf{U} = \mathbf{W}\beta + \epsilon$ for $\hat{\mathbf{U}}$

$$\hat{\mathbf{U}} = \mathbf{W}(\mathbf{W}'\mathbf{W})^{-1}\mathbf{W}'\mathbf{Y}$$

```
Uhat <- W %*% ginv(t(W) %*% W) %*% t(W) %*% U
lvector(Uhat)
```

$$\hat{\mathbf{U}} = \begin{pmatrix} 1.90 \\ 0.63 \\ 1.33 \\ 6.00 \\ 3.20 \\ 1.07 \end{pmatrix}$$

Now, we can solve for $E(\hat{\mathbf{Y}}) = \hat{\mathbf{Y}} = \mathbf{V}^{1/2}\hat{\mathbf{U}}$:

```
Yhat <- (V^{1/2})%*%Uhat
lvector(Yhat)
```

$$E(\hat{\mathbf{Y}}) = \begin{pmatrix} 1.90 \\ 1.90 \\ 4.00 \\ 6.00 \\ 3.20 \\ 3.20 \end{pmatrix}$$

(a).2 Solving for $\hat{\mathbf{C}}\beta$ given by:

$$\begin{pmatrix} 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 1 \end{pmatrix} \beta$$

We have $\hat{\mathbf{C}}\beta = \mathbf{C}(\mathbf{W}'\mathbf{W})^{-1}\mathbf{W}'\mathbf{U}$

```
CT <- matrix(c(1,1,0,0,0,1,0,1,0,0,1,0,0,1,0,1,0,0,0,1), nrow=4,ncol=5, byrow=T)
CTBetahat <- CT %*% ginv(t(W)%*%W) %*% t(W) %*% U
lvector(CTBetahat)
```

$$\hat{\mathbf{C}}\beta = \begin{pmatrix} 1.90 \\ 4.00 \\ 6.00 \\ 3.20 \end{pmatrix} = \begin{pmatrix} \widehat{\mu + \tau_1} \\ \widehat{\mu + \tau_2} \\ \widehat{\mu + \tau_3} \\ \widehat{\mu + \tau_4} \end{pmatrix}$$

(b) \mathbf{V}_2

$$\mathbf{V} = \mathbf{V}_2 = \begin{pmatrix} 1 & 1 & 0 & 0 & 0 & 0 \\ 1 & 9 & 0 & 0 & 0 & 0 \\ 0 & 0 & 9 & -1 & 0 & 0 \\ 0 & 0 & -1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & -1 \\ 0 & 0 & 0 & 0 & -1 & 9 \end{pmatrix}$$

(b).1 Transformation to OLS using $\mathbf{V}^{-1/2}$

The a square root matrix $\mathbf{V}_2^{1/2}$ can be found as follows:

- Write a matrix $\mathbf{D}^{1/2}$ containing the square root's of \mathbf{V}_2 's eigenvalues in its' diagonal.
- Write a matrix \mathbf{C} whose columns are the normalized eigenvectors of \mathbf{V}_2 .
- $\mathbf{V}_2^{1/2}$ is the product $\mathbf{C}\mathbf{D}^{1/2}\mathbf{C}'$.

We will use the square root as above to transform the GLS equations into OLS as above.

```

V <- diag(c(1,9,9,1,1,9))
V[1,2] <- 1
V[2,1] <- 1
V[4,3] <- -1
V[3,4] <- -1
V[6,5] <- -1
V[5,6] <- -1

Dh <- diag(sqrt(eigen(V)$values))
C <- eigen(V)$vectors
Vh <- C %*% Dh %*% t(C)
Vhi <- solve(Vh)

Y <- matrix(c(2, 1, 4, 6, 3, 5), nrow=6, ncol=1)
Yp <- Vhi %*% Y

X <- matrix(c(rep(1,6),
              1,1,0,0,0,0,
              0,0,1,0,0,0,
              0,0,0,1,0,0,
              0,0,0,0,1,1),nrow = 6,byrow=FALSE)
Xp <- Vhi %*% X

library(MASS)
Yphat <- Xp %*% ginv(t(Xp) %*% Xp) %*% t(Xp) %*% Yp
Yhat <- (Vh)%*%Yphat
lvector(Yhat)

CT <- matrix(c(1,1,0,0,0,1,0,1,0,0,1,0,0,1,0,1,0,0,0,1), nrow=4,ncol=5, byrow=T)
CTBetahat <- CT %*% ginv(t(Xp)%*%Xp) %*% t(Xp) %*% Yp
lvector(CTBetahat)

```

Results obtained for $\widehat{E(\mathbf{Y})}$ and $\widehat{C\beta}$ are:

$$\widehat{E(\mathbf{Y})} = \begin{pmatrix} 2.00 \\ 2.00 \\ 4.00 \\ 6.00 \\ 3.33 \\ 3.33 \end{pmatrix} \text{ and } \widehat{C\beta} = \begin{pmatrix} 2.00 \\ 4.00 \\ 6.00 \\ 3.33 \end{pmatrix}$$

3 The lm function in R allows one to do weighted least squares

with the form $\sum w_i (y_i - \hat{y}_i)^2$ for positive weights w_i . For \mathbf{V}_1 in the last question, find the BLUEs of the 4 cell means using lm and an appropriate vector of weights.

I call lm to model \mathbf{Y} as a function of the design matrix \mathbf{X} with “no intercept” and with the weights from the diagonal of \mathbf{V}_1^{-1} .

```

V <- diag(c(1,9,9,1,1,9))
Y <- matrix(c(2, 1, 4, 6, 3, 5), nrow=6, ncol=1)

```

```
X <- matrix(c(rep(1,6),
              1,1,0,0,0,0,
              0,0,1,0,0,0,
              0,0,0,1,0,0,
              0,0,0,0,1,1),nrow = 6,byrow=FALSE)
lmod <- lm(Y ~ 0 + X[,2:5], weights=diag(solve(V)))
xtable(summary(lmod))
```

Results of calling lm with weights from V_1^{-1}					
Parameter	X-location	Estimate	Std. Error	t value	Pr(> t)
$\widehat{\mu + \tau_1}$	X[, 2:5]1	1.9000	0.4743	4.01	0.0570
$\widehat{\mu + \tau_2}$	X[, 2:5]2	4.0000	1.5000	2.67	0.1165
$\widehat{\mu + \tau_3}$	X[, 2:5]3	6.0000	0.5000	12.00	0.0069
$\widehat{\mu + \tau_4}$	X[, 2:5]4	3.2000	0.4743	6.75	0.0213

The parameter estimates generated here by `lm` are identical to those obtained in question 2a, using V_1 as a source of weights.

4 By running

```
library(MASS)
data(Boston)
```

will load the Boston housing data into R. Use `?Boston` to see the information on the variables. Now create two matrixes **Y** and **X** that will be used to fit a regression model to some of these data.

Information on the variables:

```
?Boston
```

```
Y=as.matrix(Boston$medv)
X=as.matrix(Boston[,c('crim','nox','rm','age','dis')])
X=cbind(rep(1,dim(Boston)[1]),X)
```

(a) Make a scatterplot matrix for y, x_1, \dots, x_5 .

If you had to guess based on this plot, which single predictor do you think is probably the best predictor of Price? Do you see any evidence of multicollinearity (correlation among the predictors) in this graphic?

```
myscatter <- data.frame(cbind(Y,X[,c(2:6)]))
plot(myscatter)
```

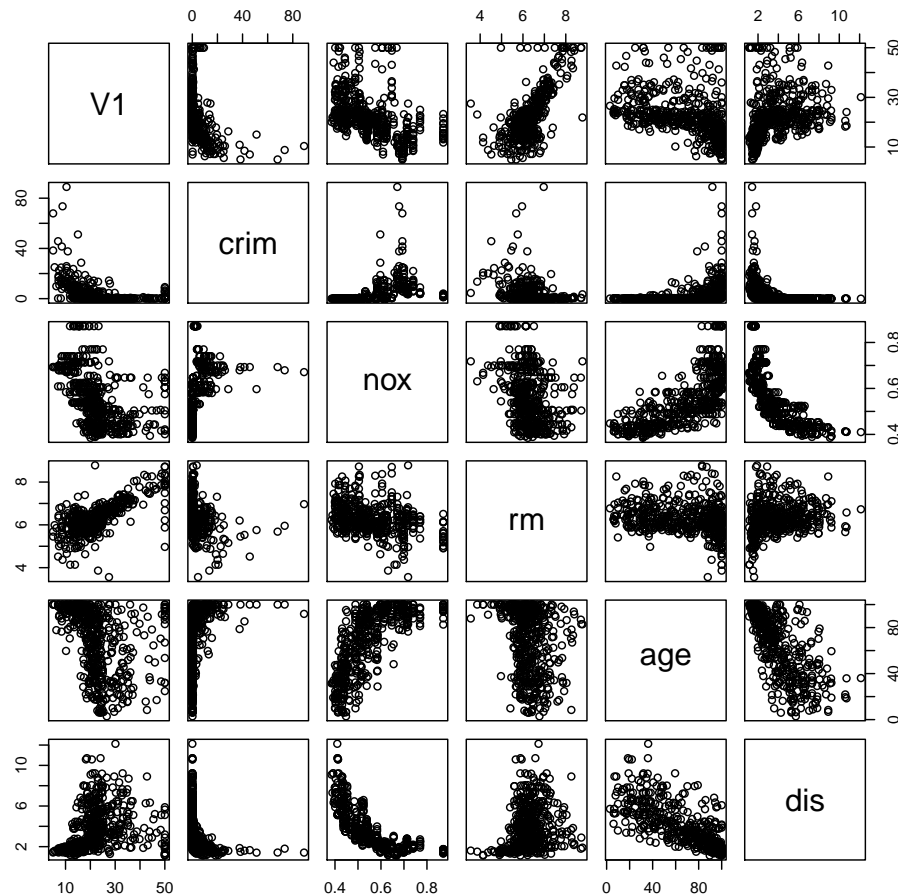


Figure 1: Scatterplot matrix showing relationships among Boston housing data variables.

Based on this scatterplot, I think that 'rm' the average number of rooms per dwelling is the best predictor of price, 'V1'. There is strong evidence of multicollinearity in the scatter plot. The 'age' or fraction of owner-occupied units built prior to 1940 and 'dis' the weighted mean of distances to five Boston employment centres appear to have a strong linear relationship, as 'age' increases, 'dis' tends to decrease. 'nox', the concentration of Nitrous Oxide has a linear relationship with 'age' as well as 'dis'.

(b) Use `qr()` function to find the rank of X

```
qr(X)$rank
```

The rank of X is 6.

(c) Use R matrix operations on the X matrix and Y vector

to find the estimated coefficient vector $\hat{\beta}$, the estimated mean vector \hat{Y} , and the vector of residuals $\mathbf{e} = Y - \hat{Y}$.

$$\hat{\beta} = (X'X)^{-1}X'Y$$

```
betahat <- ginv(t(X)%*%X) %*% t(X) %*% Y
lvector(betahat)
```


$$\hat{\beta} = \begin{pmatrix} -6.23 \\ -0.21 \\ -18.05 \\ 7.74 \\ -0.07 \\ -1.19 \end{pmatrix}$$

$$\hat{\mathbf{Y}} = \mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Y}$$

Also, $\hat{\mathbf{Y}} = \mathbf{X}\hat{\beta}$.

```
Yhat <- X %*% betahat
err <- Y- Yhat
lvector(Yhat)
lvector(err)
```

$$\hat{\mathbf{Y}} = \begin{pmatrix} 25.70 \\ 23.80 \\ 30.89 \\ \vdots \\ 28.73 \\ 27.17 \\ 21.70 \end{pmatrix}, \text{ and } \mathbf{e} = \begin{pmatrix} -1.70 \\ -2.20 \\ 3.81 \\ \vdots \\ -4.83 \\ -5.17 \\ -9.80 \end{pmatrix}$$

(d) Plot the residuals against the fitted means

```
plot(Yhat, err, xlab="Yhat, Fitted Means", ylab="Y-Yhat, Residuals",
     col= ifelse(Yhat+err >= 50, "red", "black"))
```




Figure 2: Residuals versus fitted means.

In this plot, there does seem to be non-constant variance, heteroscedasticity, which is perhaps nonlinear. The residuals tend to be farthest from zero for the smallest and largest \hat{Y} values.

There is also an artifact of the data creating a straight line in upper right hand corner of the graph. Further investigation revealed that these data points, shown in red, are points where the median value Y response value was censored at \$50,000.

(e) Create a normal plot from the values in the residuals vector.

This plot assesses the residuals for normality. I used the `qqnorm` and `qqline` functions. Here the error/residual quantiles are plotted against the theoretical quantiles for the $\text{normal}(0, \sigma^2)$ assumption.

```
qqnorm(err, ylab="Residuals", main="Normal Q-Q plot for Boston Housing Data Residuals")
qqline(err)
```

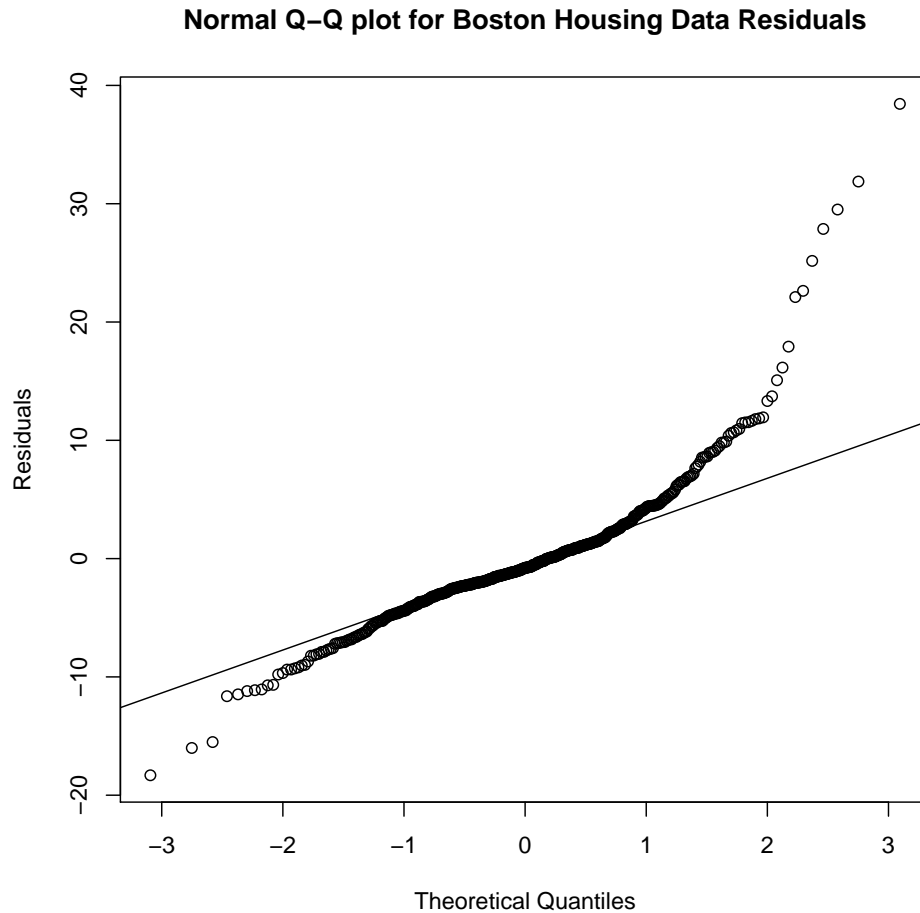


Figure 3: Normal Q-Q Plot for Boston Housing data.

While the Q-Q plot is linear for the central values, the largest and smallest residuals deviate significantly. This suggests that our underlying assumption that the errors are normally distributed may be incorrect.

(f) Compute the sum of squared residuals and the corresponding estimates of σ^2

$$\hat{\sigma}^2 = \frac{(\mathbf{Y} - \hat{\mathbf{Y}})'(\mathbf{Y} - \hat{\mathbf{Y}})}{n - \text{rank}(\mathbf{X})}$$

```
sigsqhat <- t(err) %*% err / (dim(X)[1] - qr(X)$rank)
sigsqhat
```

The estimate obtained for $\hat{\sigma}^2$ was 34.82387.

(g) Call the `lm` function in R and confirm your answers,

and note that `?lm` gives you various information such as the outputs of the function.

```
ml = lm(medv~crim+nox+rm+age+dis, data=Boston)
summary(ml)
xtable(summary(ml))
```

Residuals:				
Min	1Q	Median	3Q	Max
-18.313	-2.917	-0.785	1.979	38.442

Coefficients:				
	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-6.2273	4.0147	-1.55	0.1215
crim	-0.2081	0.0340	-6.11	0.0000
nox	-18.0509	3.9471	-4.57	0.0000
rm	7.7353	0.3954	19.56	0.0000
age	-0.0666	0.0151	-4.40	0.0000
dis	-1.1910	0.2167	-5.50	0.0000

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Call: `lm(formula = medv ~ crim + nox + rm + age + dis, data = Boston)`

Residual standard error: 5.901 on 500 degrees of freedom Multiple R-squared: 0.5924, Adjusted R-squared: 0.5883
F-statistic: 145.3 on 5 and 500 DF, p-value: < 2.2e-16

The coefficient estimates listed above are the same as those calculated earlier for `betahat`. Additionally the residual and fitted means are the same:

```
head(ml$residuals - err)
```

```
      [,1]
1  1.005196e-12
2 -1.203926e-12
3  4.233502e-12
4  5.104361e-12
5  7.077894e-12
6  1.176836e-12
```

```
head(ml$fit - Yhat)
```

```
      [,1]
1 -1.005418e-12
2  1.204370e-12
3 -4.234835e-12
4 -5.105250e-12
5 -7.077006e-12
6 -1.175948e-12
```

5 Appendix: Tangled R code

```
library(xtable)
lvector <- function(x, dig = 2, dsply=rep("f",ncol(x)+1)) {
  x <- xtable(x, align=rep("",ncol(x)+1),display=dsply,digits=dig) # We repeat empty string 6 times
  print(x, floating=FALSE, tabular.environment="pmatrix",
        hline.after=NULL, include.rownames=FALSE, include.colnames=FALSE)
}
```

```

V <- diag(c(1,9,9,1,1,9))
Vhi <- solve(V^(1/2))
lvector(Vhi)

Y <- matrix(c(2, 1, 4, 6, 3, 5), nrow=6, ncol=1)
U <- Vhi %*% Y
lvector(U)

X <- matrix(c(rep(1,6),
               1,1,0,0,0,0,
               0,0,1,0,0,0,
               0,0,0,1,0,0,
               0,0,0,0,1,1),nrow = 6,byrow=FALSE)
W <- Vhi %*% X
lvector(W)

Uhat <- W %*% ginv(t(W) %*% W) %*% t(W) %*% U
lvector(Uhat)

Yhat <- (V^{1/2})%*%Uhat
lvector(Yhat)

CT <- matrix(c(1,1,0,0,0,1,0,1,0,0,1,0,0,1,0,1,0,0,0,1), nrow=4,ncol=5, byrow=T)
CTBetahat <- CT %*% ginv(t(W)%*%W) %*% t(W) %*% U
lvector(CTBetahat)

V <- diag(c(1,9,9,1,1,9))
V[1,2] <- 1
V[2,1] <- 1
V[4,3] <- -1
V[3,4] <- -1
V[6,5] <- -1
V[5,6] <- -1

Dh <- diag(sqrt(eigen(V)$values))
C <- eigen(V)$vectors
Vh <- C %*% Dh %*% t(C)
Vhi <- solve(Vh)

Y <- matrix(c(2, 1, 4, 6, 3, 5), nrow=6, ncol=1)
Yp <- Vhi %*% Y

X <- matrix(c(rep(1,6),
               1,1,0,0,0,0,
               0,0,1,0,0,0,
               0,0,0,1,0,0,
               0,0,0,0,1,1),nrow = 6,byrow=FALSE)
Xp <- Vhi %*% X

library(MASS)

```

```

Yphat <- Xp %*% ginv(t(Xp) %*% Xp) %*% t(Xp) %*% Yp
Yhat <- (Vh)%*%Yphat
lvector(Yhat)

CT <- matrix(c(1,1,0,0,0,1,0,1,0,0,1,0,0,1,0,1,0,0,0,1), nrow=4,ncol=5, byrow=T)
CTBetahat <- CT %*% ginv(t(Xp)%*%Xp) %*% t(Xp) %*% Yp
lvector(CTBetahat)

V <- diag(c(1,9,9,1,1,9))
Y <- matrix(c(2, 1, 4, 6, 3, 5), nrow=6, ncol=1)
X <- matrix(c(rep(1,6),
               1,1,0,0,0,0,
               0,0,1,0,0,0,
               0,0,0,1,0,0,
               0,0,0,0,1,1),nrow = 6,byrow=FALSE)
lmod <- lm(Y ~ 0 + X[,2:5], weights=diag(solve(V)))
xtable(summary(lmod))

pdf(file="HW2_4a.pdf")
myscatter <- data.frame(cbind(Y,X[,c(2:6)]))
plot(myscatter)
dev.off()

qr(X)$rank

betahat <- ginv(t(X)%*%X) %*% t(X) %*% Y
lvector(betahat)

Yhat <- X %*% betahat
err <- Y- Yhat
lvector(Yhat)
lvector(err)

sigsqhat <- t(err) %*% err / (dim(X)[1] - qr(X)$rank)
sigsqhat

ml = lm(medv~crim+nox+rm+age+dis, data=Boston)
summary(ml)
xtable(summary(ml))

head(ml$residuals - err)

head(ml$fit - Yhat)

```